



File Operation


Sujoy Saha
Department of CSE
NIT Durgapur

Types of Files in C

Types of Files in C



file.txt

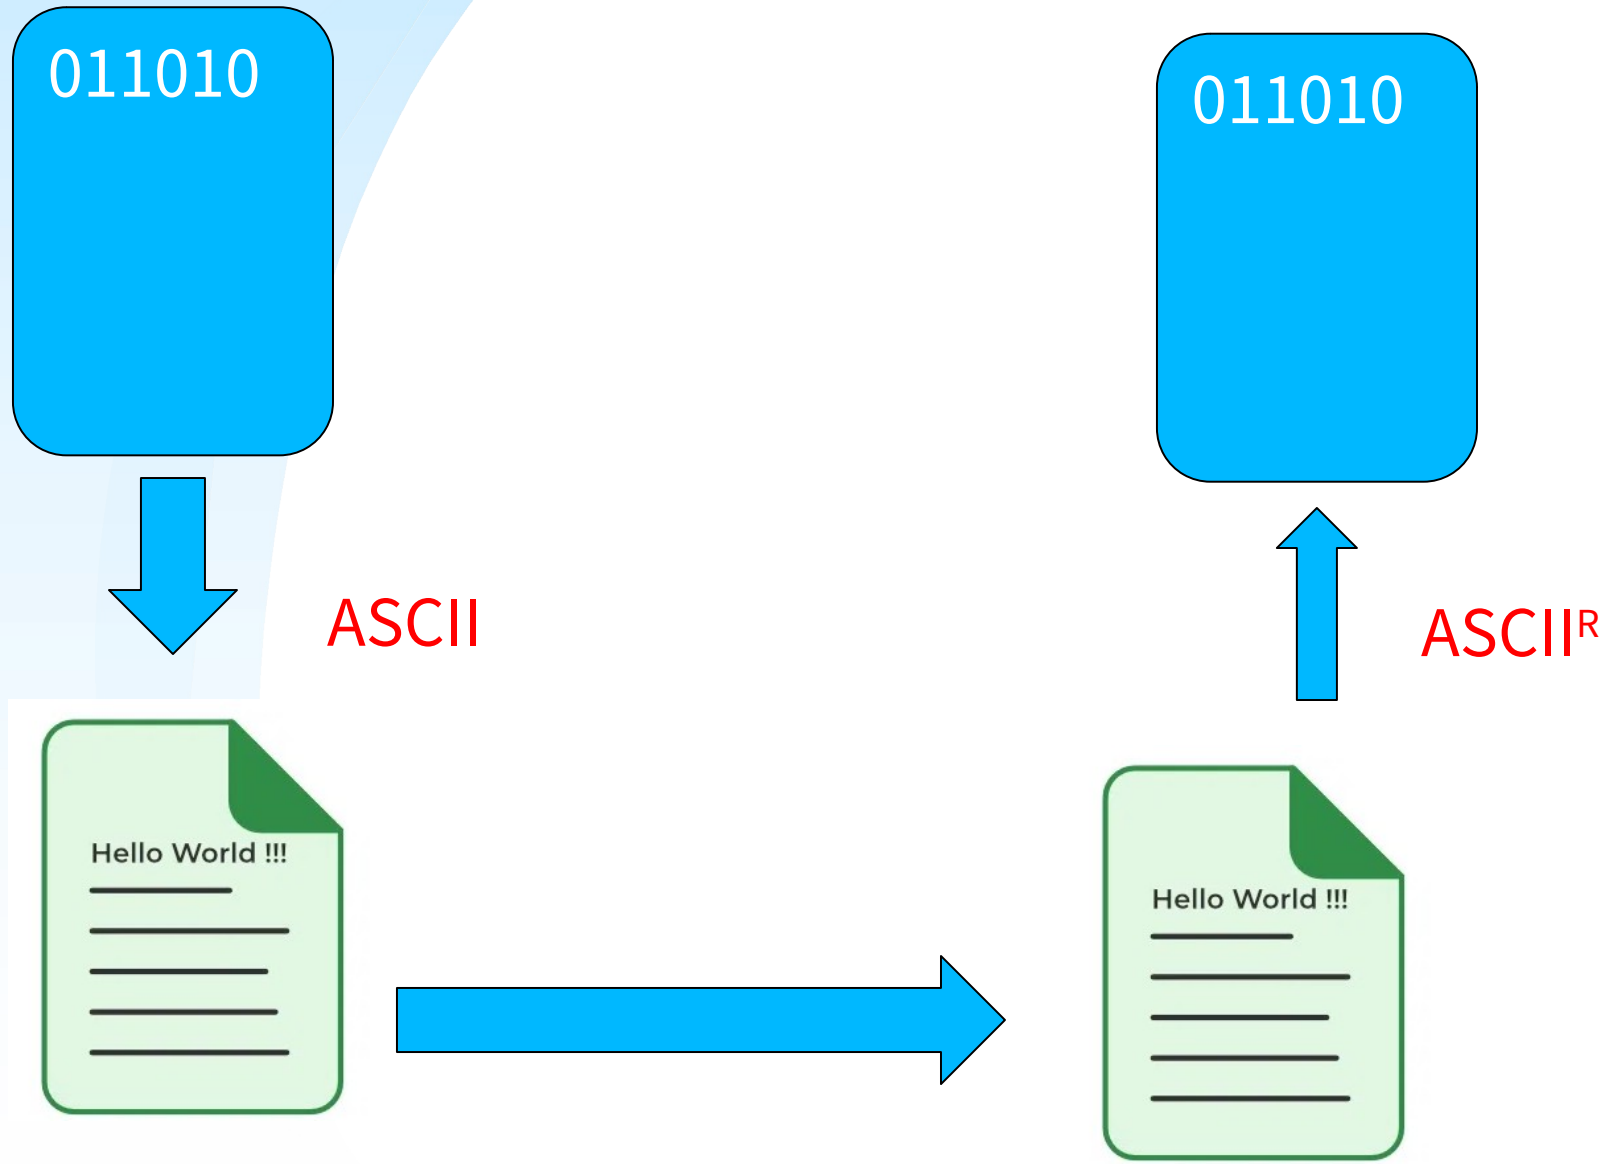


file.bin

Text Files: A text file contains data in the **form of ASCII characters** and is generally used to store a stream of characters.

Binary Files: A binary file contains data in **binary form (i.e. 0's and 1's)** instead of ASCII characters. They contain data that is stored in a similar manner to how it is stored in the main memory.

Types of Files in C






When dealing with audio files (*like .wav, .mp3, .flac, etc.*) in C, the correct choice is almost always to use binary mode (*"rb", "wb", "ab", etc.*) rather than *text/ASCII mode* (*"r", "w", "a"*).


Types of Files in C

011010



CODEC

011010



CODEC^R



Why Binary Mode is Better (and Safer) for Audio/Video Files

- ***Audio files are binary by nature***
 - They don't store human-readable text, but raw byte sequences representing samples, headers, metadata, etc.
- the file is read and written **exactly as it is, byte by byte.**
 - This ensures that the ***data is preserved without any modification***, making it ideal for non-text files like images, videos, or executables.
- ***Text mode may corrupt the file***
 - In text mode, the runtime library may automatically translate certain characters (e.g., '\n' into "\r\n" on Windows).
 - That breaks the structure of audio files, since even one altered byte can make the file unreadable.
 - There is **no risk of accidental data corruption due to character encoding or newline conversions.**

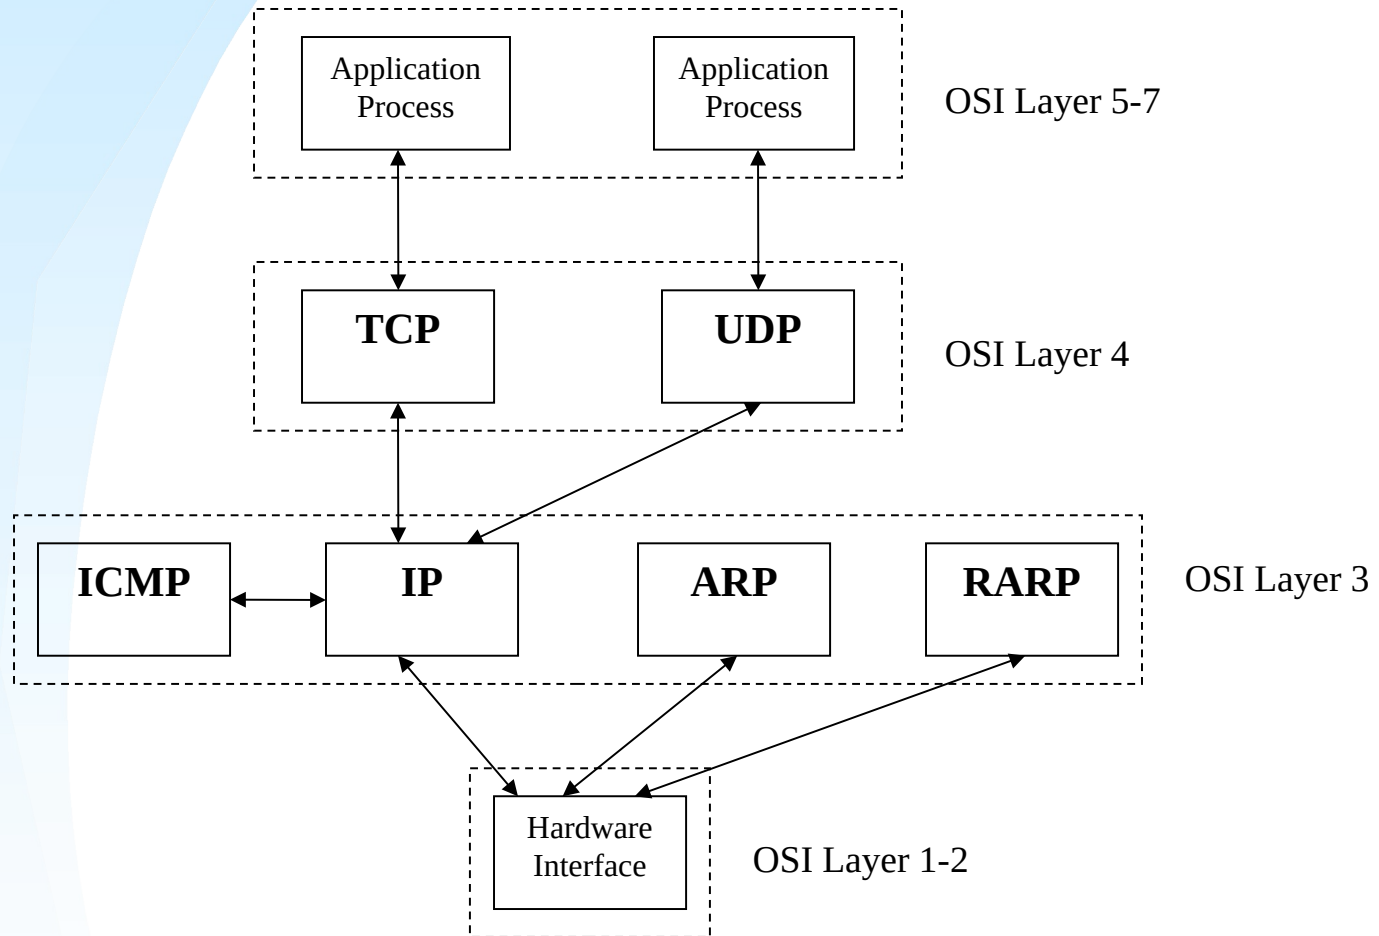
Why Binary Mode is Better (and Safer) for Audio/Video Files

- ***Portability***
 - If you write in text mode, behavior may differ between platforms (Windows vs Linux vs macOS).
 - Binary mode guarantees consistent, byte-for-byte I/O across systems.

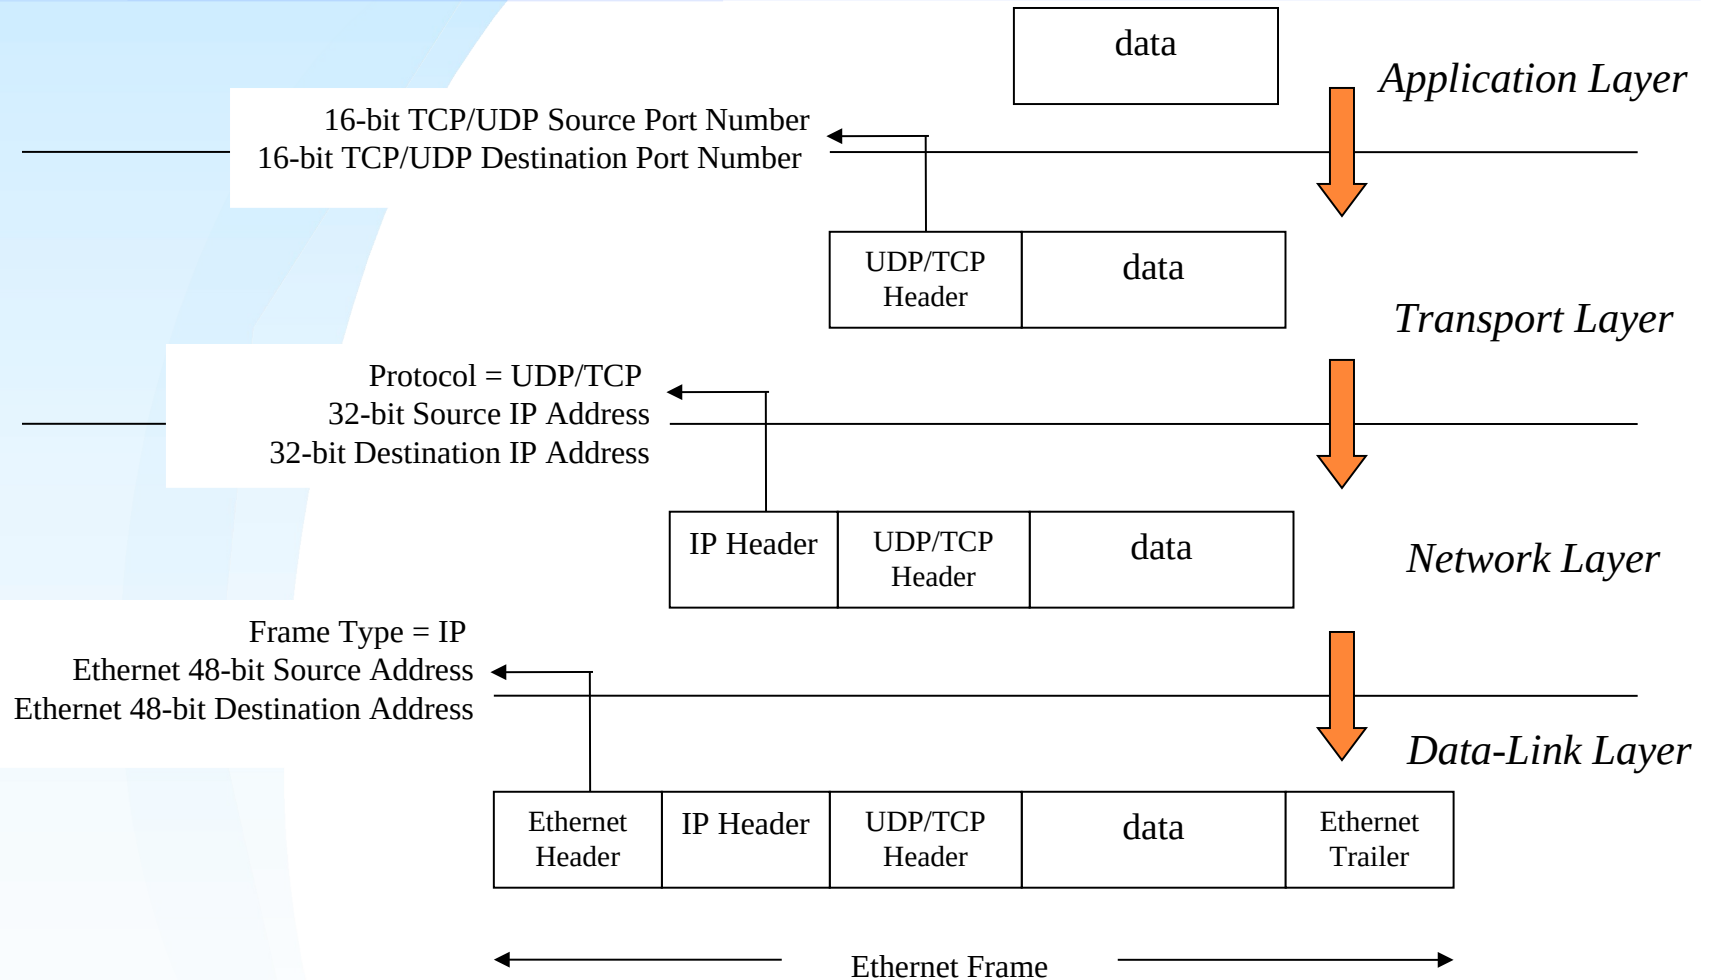
TCP/IP Socket Programming ***in C***

Sujoy Saha
Department of CSE
NIT Durgapur

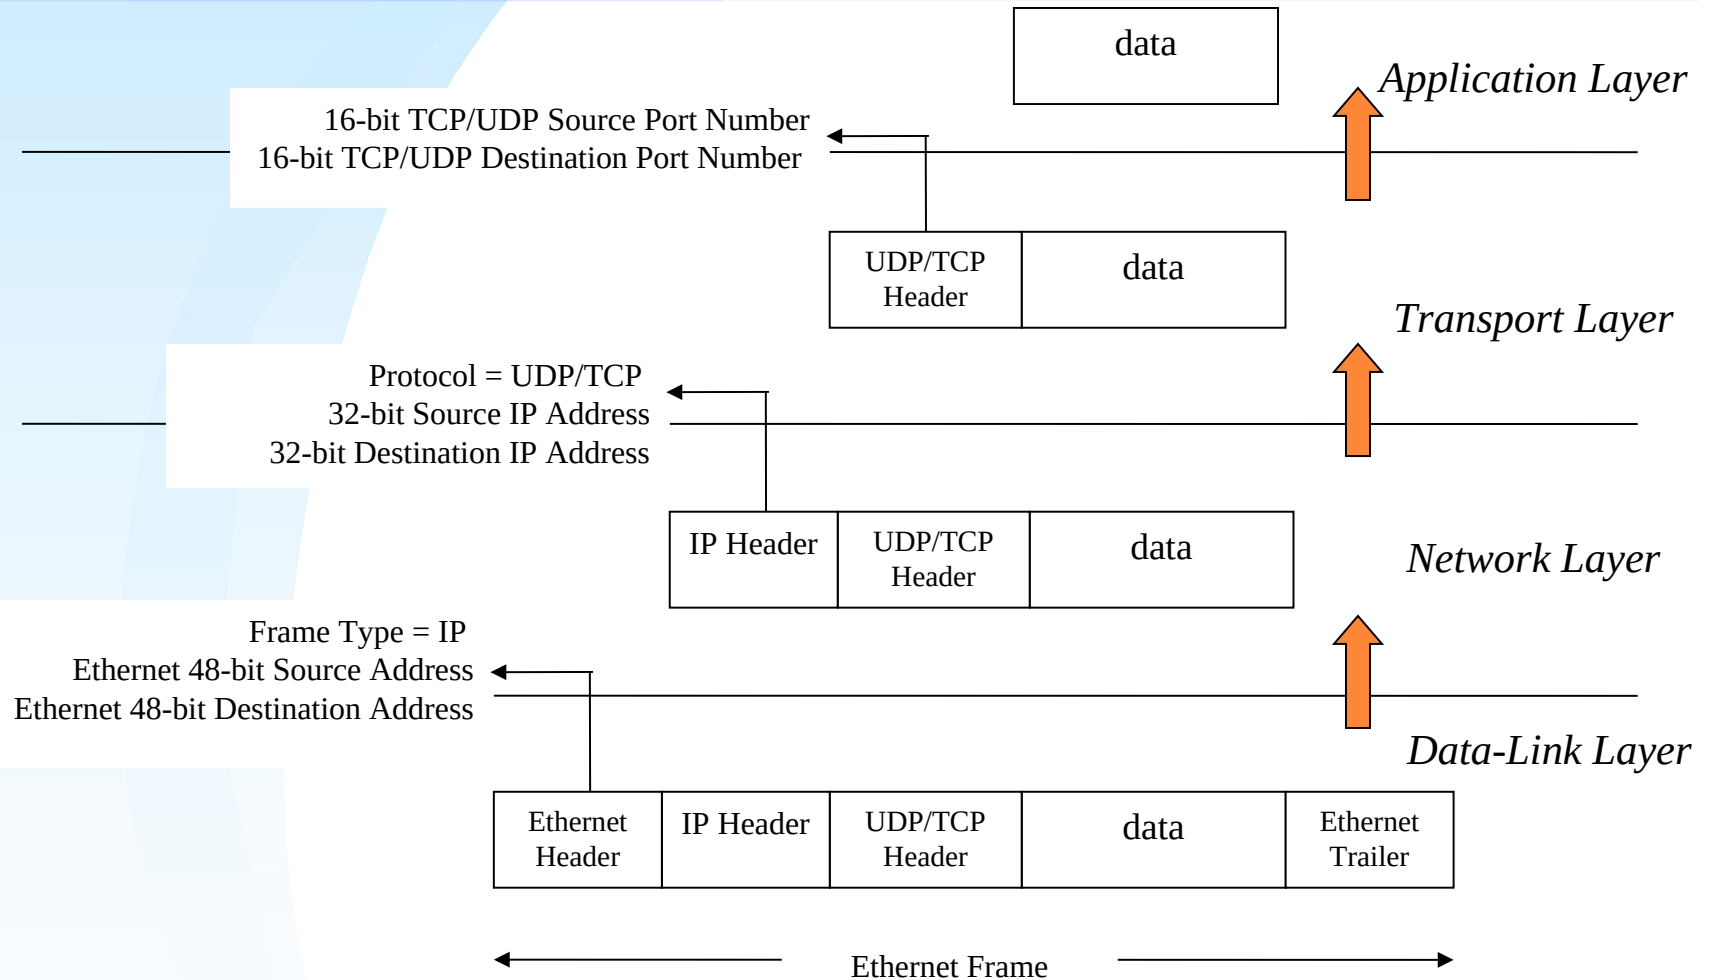
TCP/IP Protocol Suite



Sending Data




Receiving Data



Socket

“A **network socket** is an endpoint of an inter-process communication across a computer network.”



Connection

- In TCP/IP protocol suite
 - a connection defines the communication link between two processes
 - An *association* defined by the 5-tuple completely specifies two processes that make up a connection

{protocol, local-addr, local-port, foreign-addr, foreign-port}

- The protocol (TCP or UDP)
- The local host's IP address (32-bit)
- The local port number (16-bit)
- The foreign host's IP address (32-bit)
- The foreign port number (16-bit)

- Example

{tcp, 192.168.2.2, 1500, 192.168.2.10, 21}

Sockets

- Sockets are *application program interface (API)* to the communication protocol.
- Availability of an API depends both on the operating system and the programming language.

{tcp/UDP, 192.168.2.2, 1500,
192.168.2.10, 21}

Socket Addresses

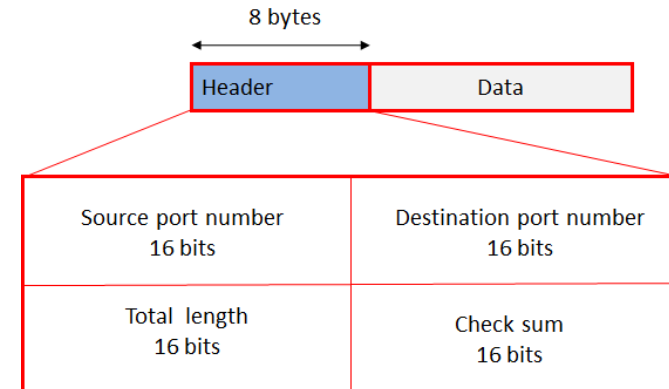
- For Internet family, following structures are defined in

`<netinet/in.h>`

```
struct in_addr {  
    u_long    s_addr;    /*32-bit netid/hostid*/  
};
```

```
struct sockaddr_in {  
    short      sin_family; /*AF_INET*/  
    u_short    sin_port;   /*16-bit port number*/  
    struct in_addr sin_addr; /*32-bit netid/hostid*/  
    char       sin_zero[8]; /*unused*/  
};
```

- The header file `<sys/types.h>` provides data type definitions



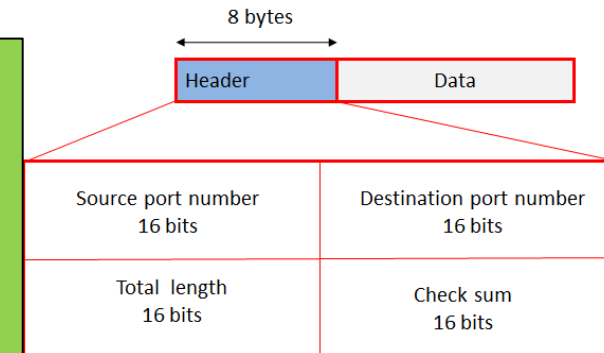
{tcp/UDP, 192.168.2.2, 1500,
192.168.2.10, 21}

Socket Addresses

- For Internet family, following structures are defined in

In the socket API, the family (sometimes called domain) tells the operating system *what kind of addresses the socket will use*.

- AF_INET → IPv4 addresses (e.g., 192.168.1.1)
- AF_INET6 → IPv6 addresses (e.g., 2001:db8::1)
- AF_UNIX / AF_LOCAL → Local (Unix domain) sockets for communication on the same machine – **Not use IP Address**
 - it uses a file **path** in the filesystem (e.g., /tmp/mysocket) to identify the communication endpoint.

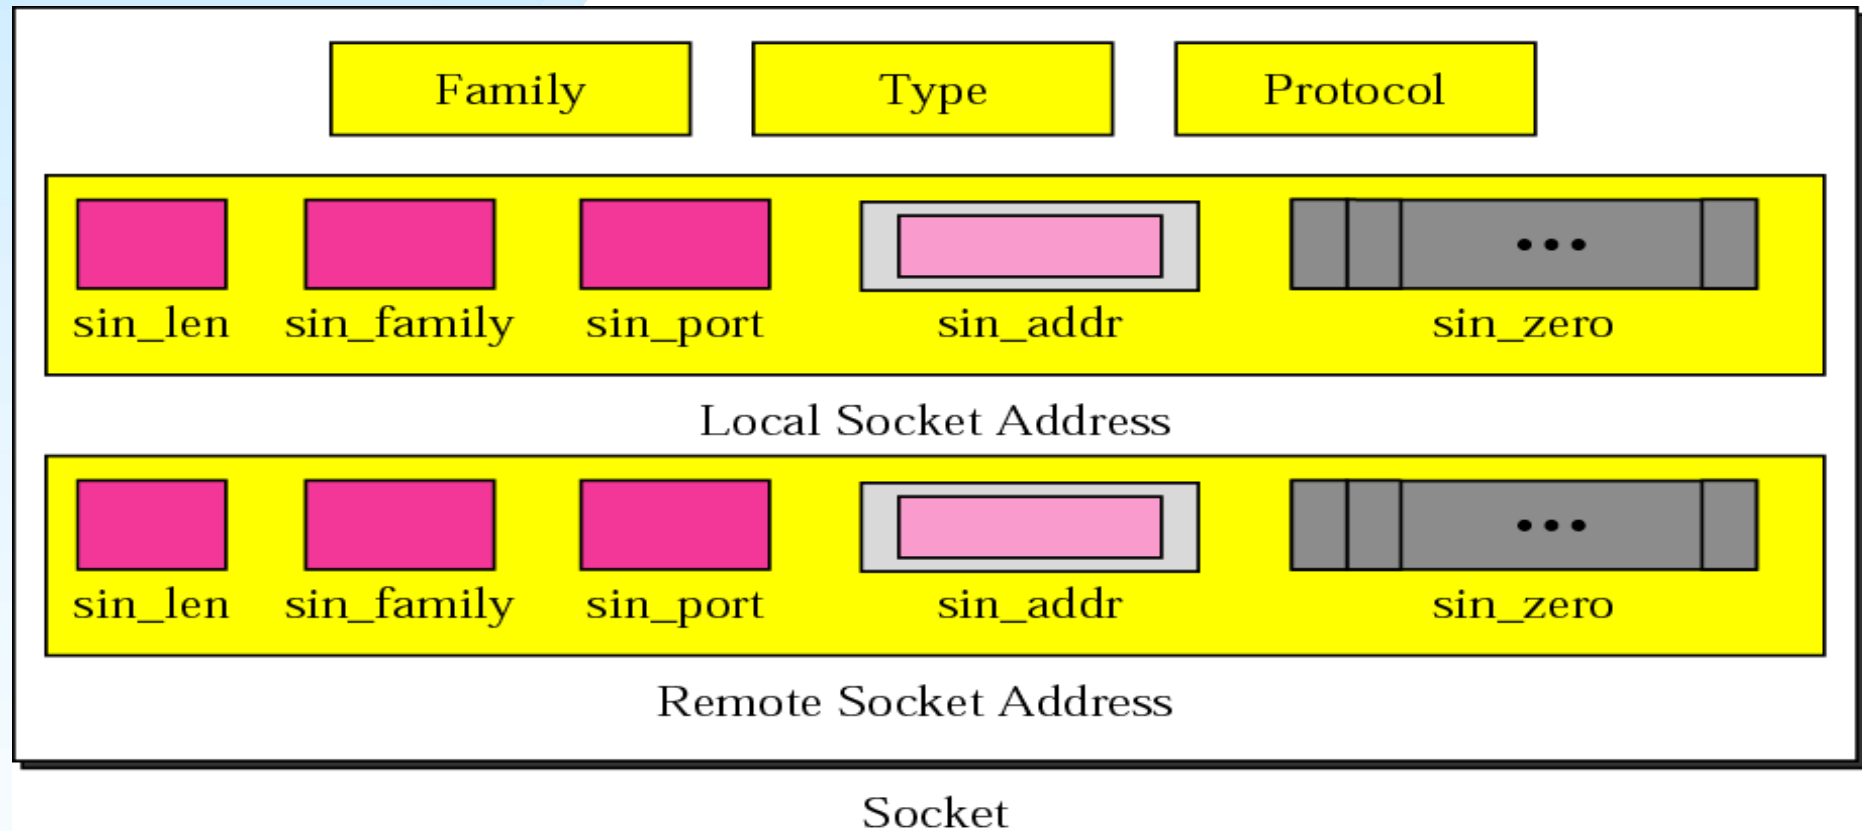


tid*/

data type

Figure 16-4

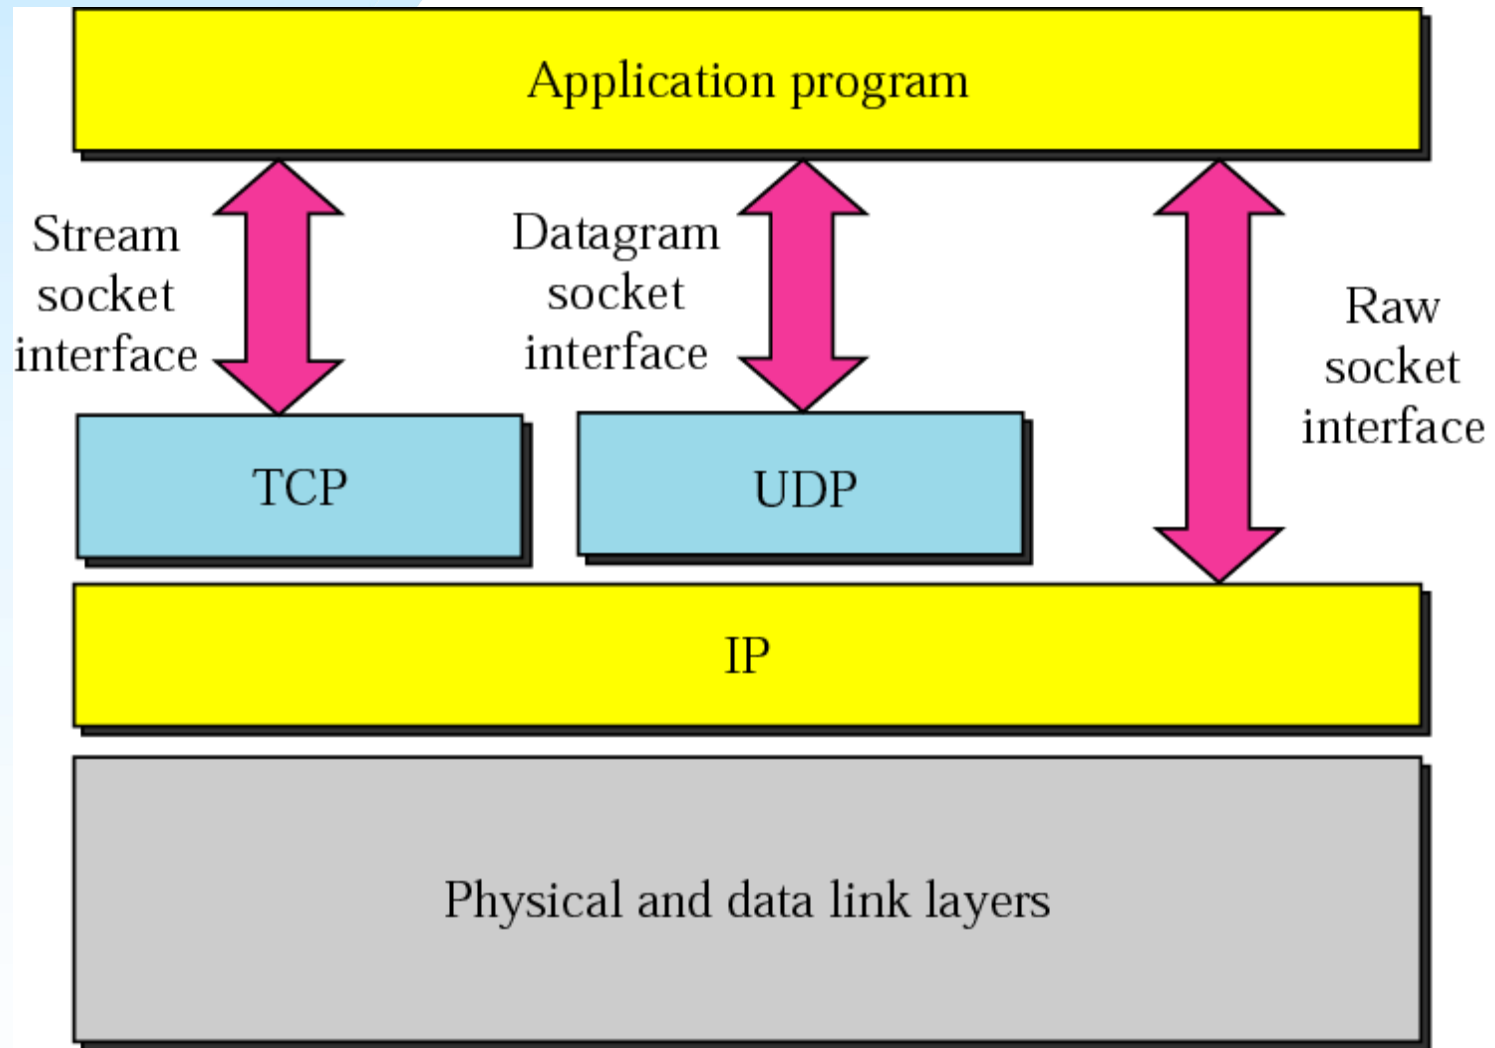
Socket structure



Communication domain in which the socket should be created. Some of address families are AF_INET (IP), AF_INET6 (IPv6), AF_UNIX (local channel, similar to pipes), AF_ISO (ISO protocols), and AF_NS (Xerox Network Systems protocols).

Figure 16-5

Socket types





BYTE ORDERING

Figure 16-6

Big-endian byte order

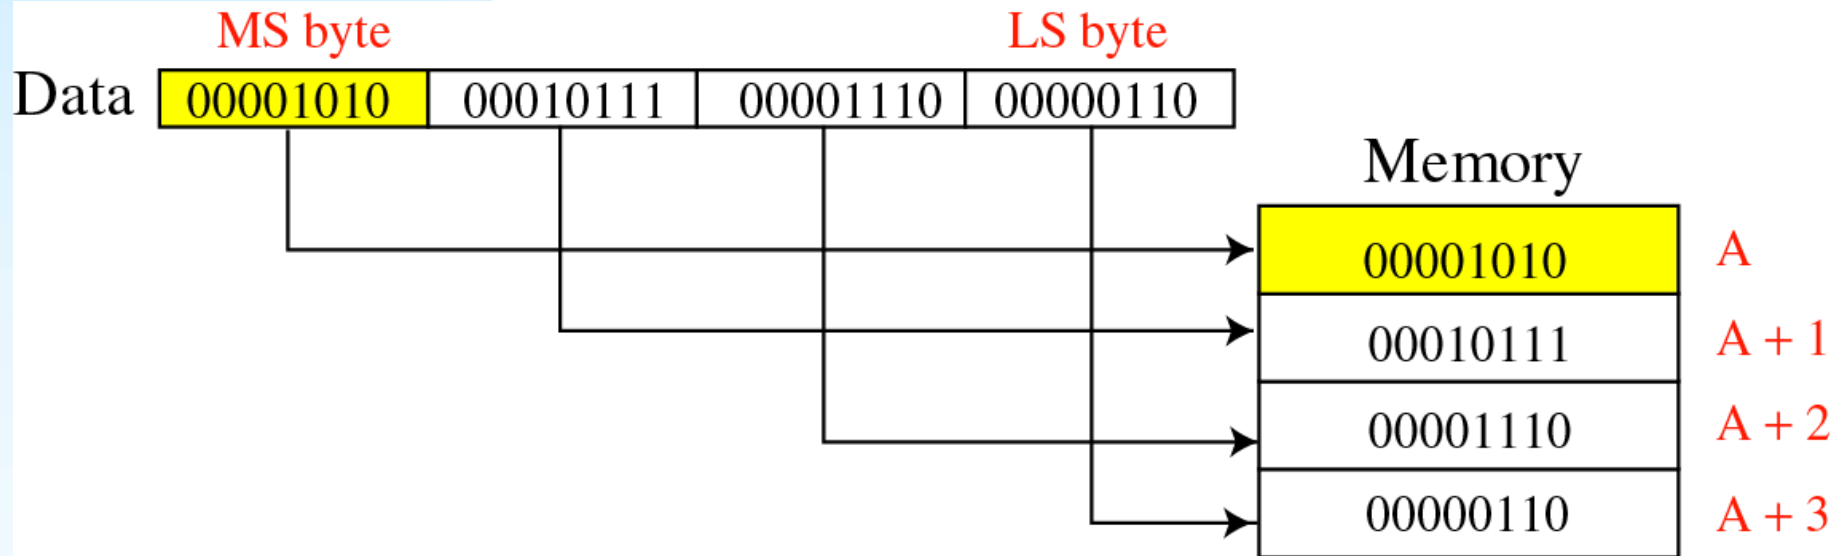
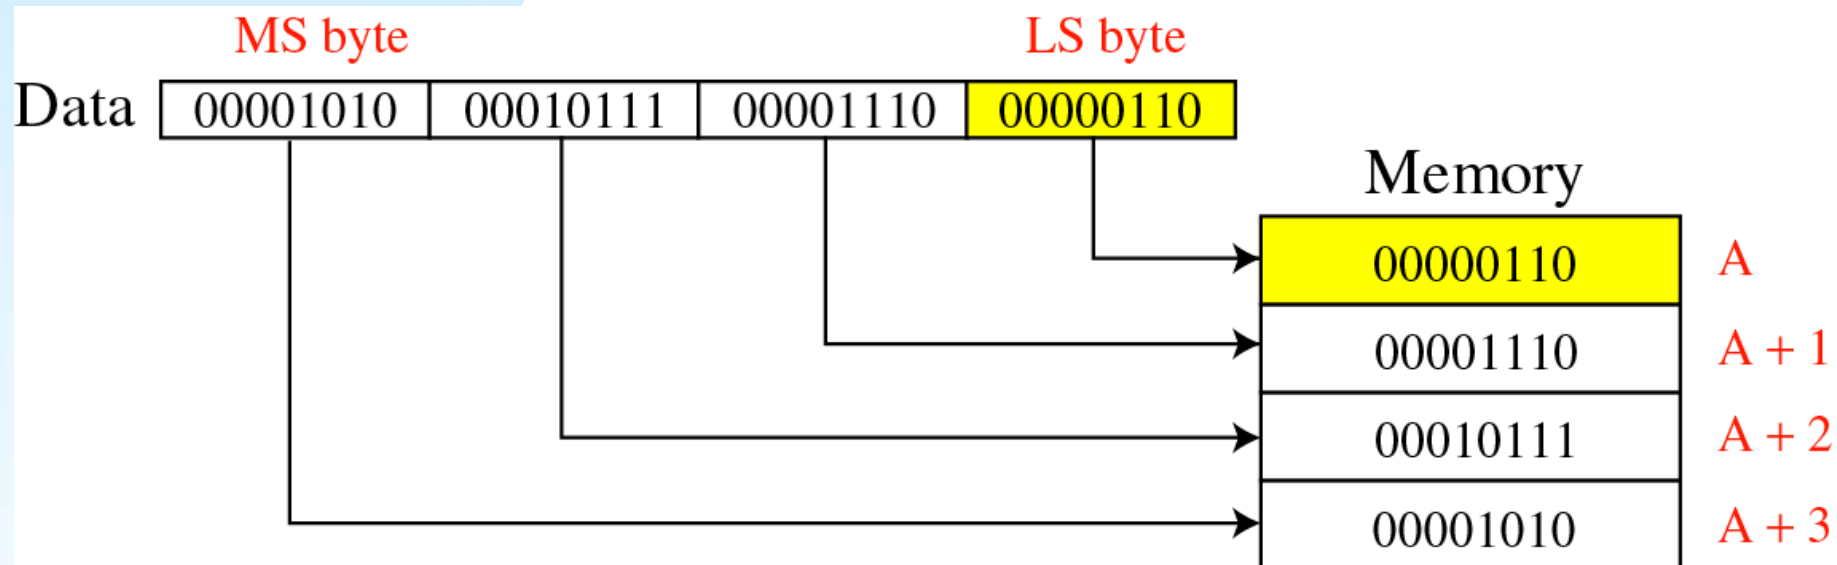


Figure 16-7

Little-endian byte order



Note

The byte order for the TCP/IP protocol suite is big endian.

Byte Ordering Functions

- Hosts can store multi-byte values differently (*host byte order*)
 - Little-endian byte order
 - Big-endian byte order
- Network protocols must specify a *network byte order*
 - Internet protocols use big-endian byte ordering
- Byte ordering routines

```
u_long htonl(u_long hostlong);  
u_short htons(u_short hostshort);  
u_long ntohl(u_long netlong);  
u_short ntohs(u_short netshort);
```


u_long: A data type that can store **non-negative integers**. The size is usually **4 (32 bit machine)** or **8 bytes (64 bit machine)**, depending on the system.

Often defined as a shorthand for **unsigned long**, usually in **system headers or legacy code**.

It's not part of the standard C library but may appear in codebases or system-specific headers.

Figure 16-8

Bite-order transformation

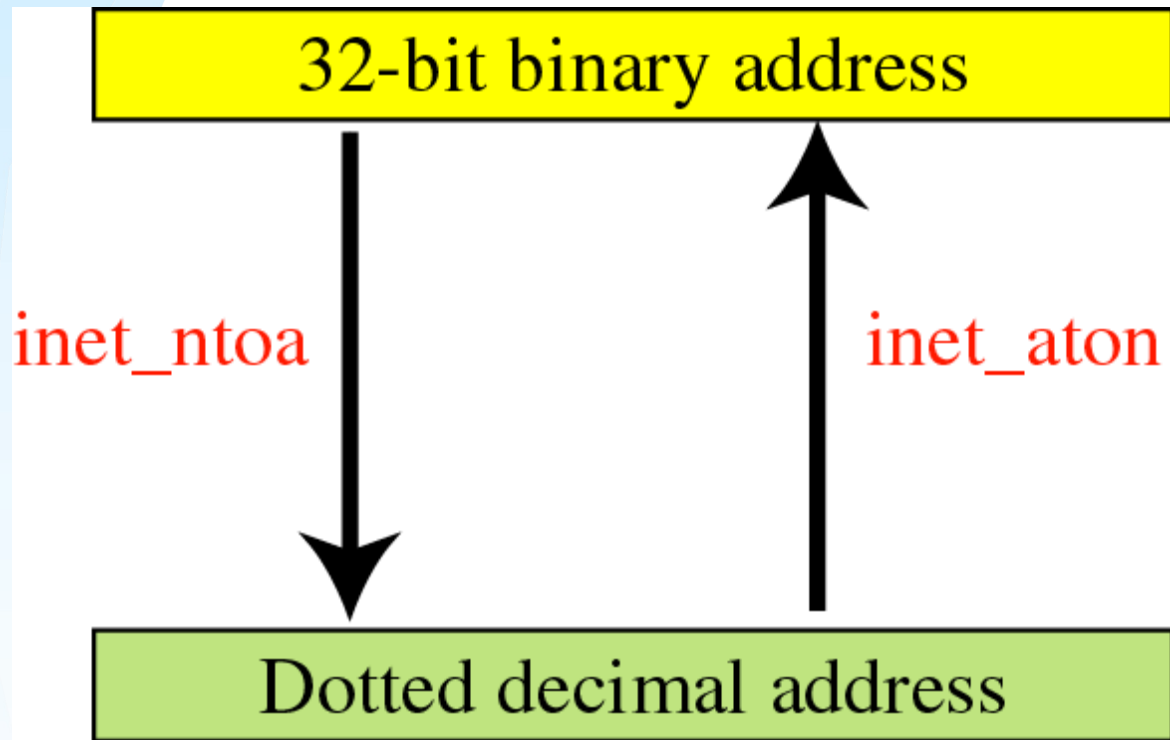


Data types

```
#include<stdint.h>
```

int8_t	Signed 8-bit integer	-128 to 127 [7 bit → number and 1 bit for sign]
int16_t	Signed 16-bit integer	
int32_t	Signed 32-bit integer	
uint8_t	Unsigned 8-bit integer	Only positive numbers→0-255[utilizing all 8 bits]
uint16_t	Unsigned 16-bit integer	
uint32_t	Unsigned 32-bit integer	

Address transformation



The `inet_ntoa` function in C is used to convert an **IP address in a numerical format** (stored as a struct `in_addr`) into a **human-readable dotted-decimal string** (e.g., 192.168.1.1).

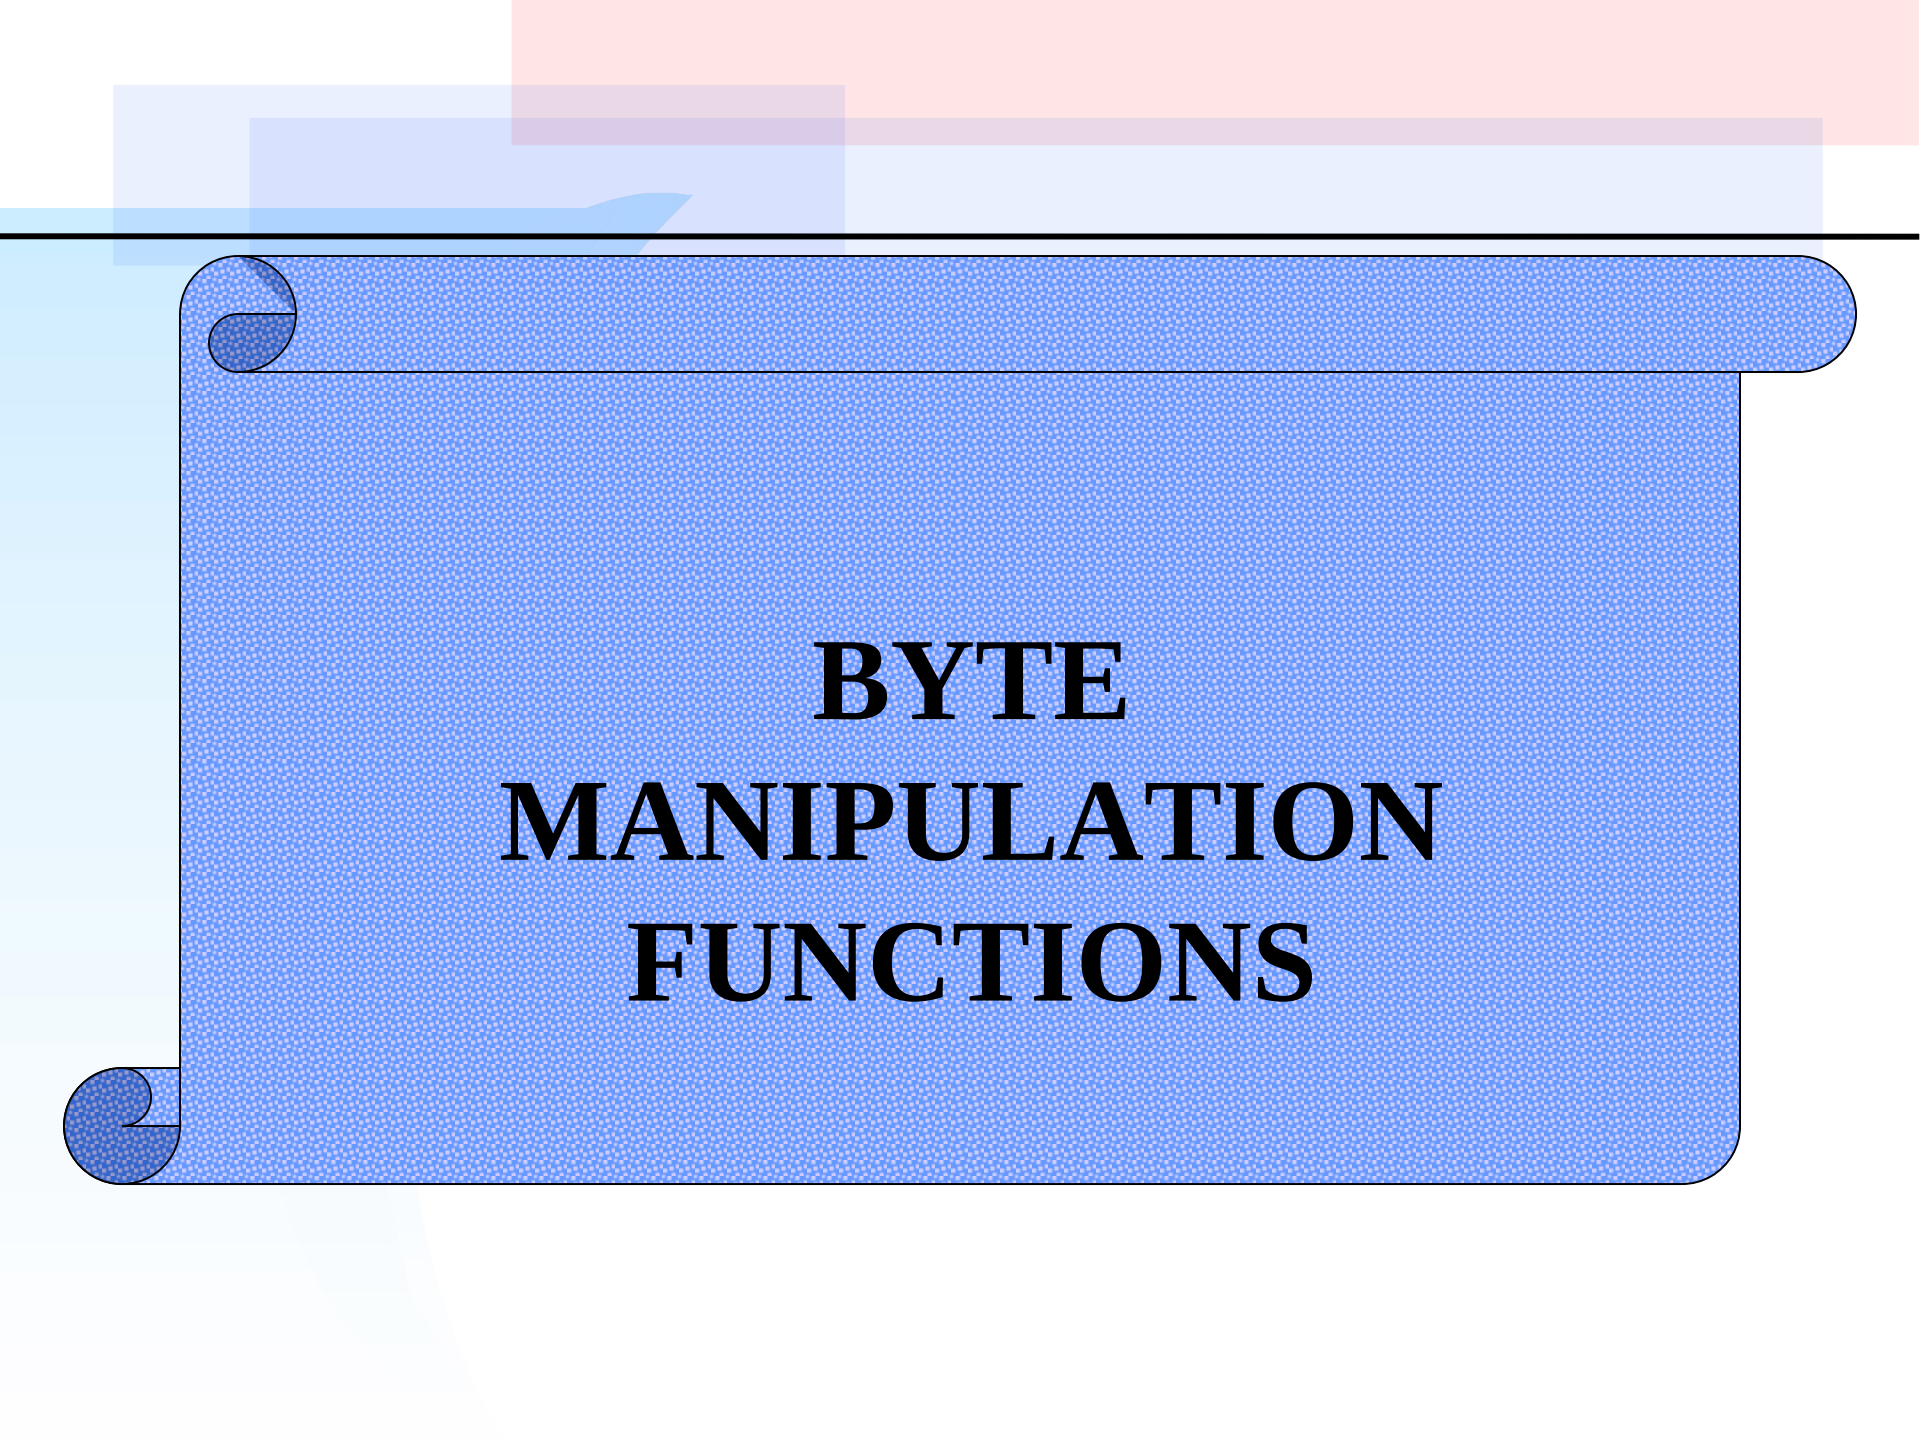
Address Conversion Routines

```
unsigned long inet_addr(char *ptr);  
int inet_aton(char *ptr, struct in_addr *addrptr);
```

Converts a character string in dotted-decimal notation to a 32-bit Internet address in network byte order

```
char *inet_ntoa(struct in_addr inaddr);
```

Converts a 32-bit Internet address in network byte order to a character string in dotted decimal notation



BYTE MANIPULATION FUNCTIONS

Declarations for byte-manipulation functions

```
void      *memset ( void *dest , int  chr , size_t len ) ;
```

```
void      *memcpy ( void *dest , const void *src , size_t len ) ;
```

```
int        memcmp ( const void *first , const void *second , size_t len ) ;
```

```
char buffer[1400];
```

```
memset(buffer,0,sizeof(buffer));
```

```
bzero(buffer, sizeof(buf));
```

```
memcpy(buffer, msg.m_txt, sizeof(msg.m_txt));
```

```
struct MSG
{
    int m_txt[MAX_TEXT];
    int m_seq;
    int
    total_no_of_packets;
};
```



SOCKET SYSTEM CALLS

socket System Call

```
int socket(int family, int type, int protocol);
```

- The family is one of

AF_UNIX	Unix internal protocols
AF_INET	Internet Protocols
AF_NS	Xerox NS Protocols
AF_IMPLINK	IMP link Layer

- AF_INET uses the TCP/IP protocol.
- AF_UNIX creates filesystem objects and it only works between processes on the same host.
- AF_NS is a set of protocols that were used by Xerox Systems for data communication. Its basic **working mechanism** is almost the **same as in the TCP/IP protocol** suit, but **XNS contains only two network layers**

AF_INET and AF_UNIX

- AF_INET uses the TCP/IP protocol.
- AF_UNIX creates filesystem objects and it only works between processes on the same host.
- AF_UNIX is much faster than AF_INET.
- Xerox Network Systems (XNS) is a set of protocols that were used by Xerox Systems for data communication. Xerox used **XNS for file transfers, sharing network resources, packet transfers, sharing routing information and remote procedure calls**. Its basic **working mechanism** is almost the **same as in the TCP/IP protocol** suit, but **XNS contains only two network layers**. This differs from the seven-layer Open Systems Interconnection (OSI) model, although the functionality is basically the same.

socket System Call

```
int socket(int family, int type, int protocol);
```

- The family is one of

AF_UNIX Unix internal protocols

AF_INET Internet Protocols

AF_NS Xerox NS Protocols

AF_IMPLINK IMP link Layer

- The socket type is one of the following

SOCK_STREAM Stream Socket

SOCK_DGRAM Datagram Socket

SOCK_RAW Raw Sockets

SOCK_SEQPACKET Sequenced Packet Socket

SOCK_RDM Reliably Delivered Message Socket

- The protocol argument to the socket system call is typically set to 0
- The socket system call returns an integer value called socket descriptor
- This number is passed as a parameter to almost all of the other library calls.

Figure 16-15

Declaration for *socket* function

```
int socket (int family , int type , int protocol) ;
```

↑ ↑ ↑
 AF_INET SOCK_DGRAM 0
 SOCK_STREAM

Returns a socket descriptor if successful; -1 if error.



- The `socket()` system call returns a file descriptor.
- In fact, a socket is similar to an opened file because it is possible to read and write data on it by means of the usual `read()` and `write()` system calls.
- in the context of the client-server model, the socket system calls needed to create and connect a pair of sockets and transmit data.

<i>type</i> \ <i>family</i>	AF_INET
SOCK_STREAM	TCP or SCTP
SOCK_DGRAM	UDP
SOCK_SEQPACKET	SCTP
SOCK_RAW	IPv4

<i>Protocol</i>
IPPROTO_TCP
IPPROTO_UDP
IPPROTO_SCTP



Family

AF_UNIX

AF_INET

AF_ISO

AF_NS

AF_IPX

AF_APPLETALK

Description

UNIX internal (file system sockets)

ARPA Internet protocols (UNIX network sockets)


ISO standard protocols

Xerox Network Systems protocols

Novell IPX protocol

Appletalk DDS

bind



bind system call:

- `bind()` associates the socket with its local address
- The `bind()` system call **binds an address (IP & PORT)**



```
#include <sys/socket.h>
int bind( int sockfd, const struct sockaddr * myaddr, socklen_t
addrlen );
```

Returns: 0 if OK, -1 on error

sockaddr is a generic structure, bind system call uses type punning

■ e.g. `bind()`

```
myaddr.sin_family = AF_INET;
myaddr.sin_addr.s_addr = htonl( INADDR_ANY );
myaddr.sin_port = htons( 7 ); // 7 for echo service
```

```
if( bind( sockfd, (struct sockaddr *)&myaddr,
sizeof(myaddr) ) < 0 ) { ... }
```

bind System Call

- Uses of bind system call
 - Servers register their well-known address with the system; both connection-oriented and connection-less servers need to do this
 - A client can register a specific address for itself
- The bind system call specifies the *local-addr* and *local-port* elements of the 5-tuple of an association

```
struct sockaddr_in {  
    short    sin_family; /*AF_INET*/  
    u_short  sin_port;   /*16-bit port number*/  
    struct in_addr sin_addr; /*32-bit netid/hostid*/  
    char sin_zero[8]; /*unused*/  
};
```

```
Struct sockaddr  
{  
    unsigned short sa_family;  
    Char sa_data[14]  
}
```

$\text{sin_port (2 Byte) + Sin_addr(4 Byte) + sin_Zero (8Byte) = 14 byte}$

Why is a bind system call required in socket programming?

- Otherwise if a port is not binded to any socket descriptor, we can use the same port and IP combination again to connect to a different destination.
- `bind()` assigns the address specified to by `addr` to the socket referred to by the file descriptor `sockfd`.
- `bind()` associates the socket with its local address [that's why server side binds, so that clients can use that address to connect to server.]

sendto and recvfrom Functions

```
int sendto(int sockfd, char *buff, int length, int flags,  
struct sockaddr *to, socklen_t addrlen);
```

to argument is a socket address structure containing protocol address of where the data is to be sent

sockfd -- The socket descriptor.

buff --The pointer to the buffer containing the message to transmit.

length -- The length of the message in the buffer pointed to by the msg parameter.

Flags-- **MSG_DONTROUTE** -- Don't use a gateway to send out the packet, only send to hosts on directly connected networks. This is usually used only by diagnostic or routing programs ;

MSG_DONTWAIT - Enables nonblocking operation;

MSG_WAITALL requests blocking until the entire number of bytes requested can be read.

to --The address of the target.

addr_len --The size of the address pointed to by address.

sendto and recvfrom Functions

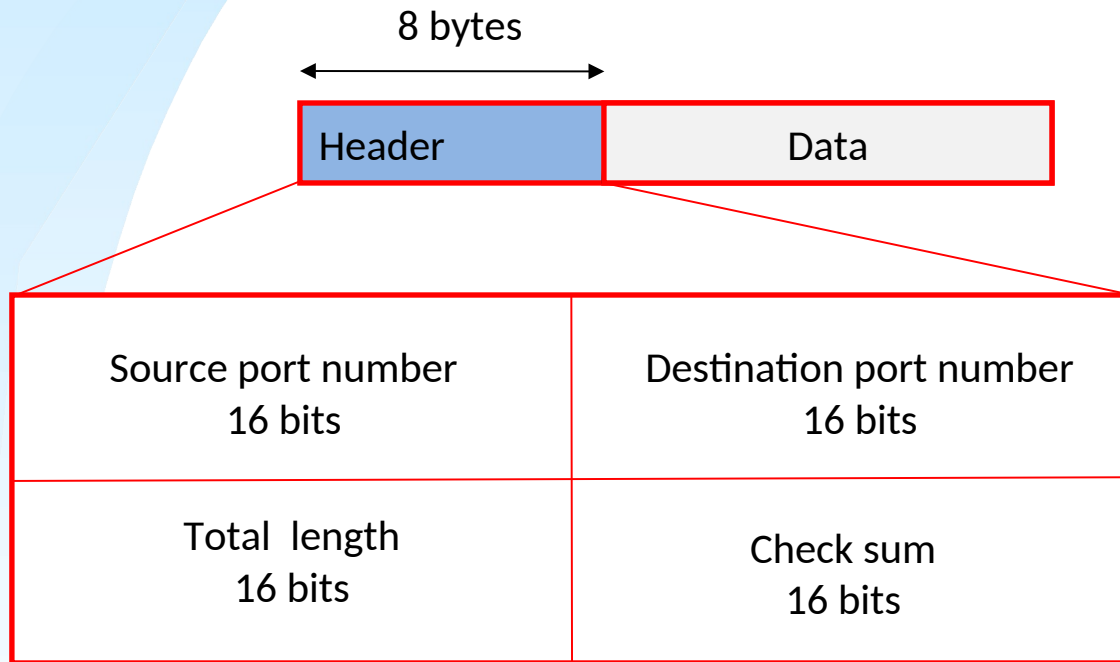
- A blocking system call is one that must wait until the action can be completed. `read()` would be a good example - if no input is ready, it'll sit there and wait until some is (provided you haven't set it to non-blocking, of course, in which case it wouldn't be a blocking system call)
- If there is not enough available buffer space to hold the socket data to be transmitted, and the socket is in blocking mode, `sendto()` blocks the caller until additional buffer space becomes available.

sendto and recvfrom Functions

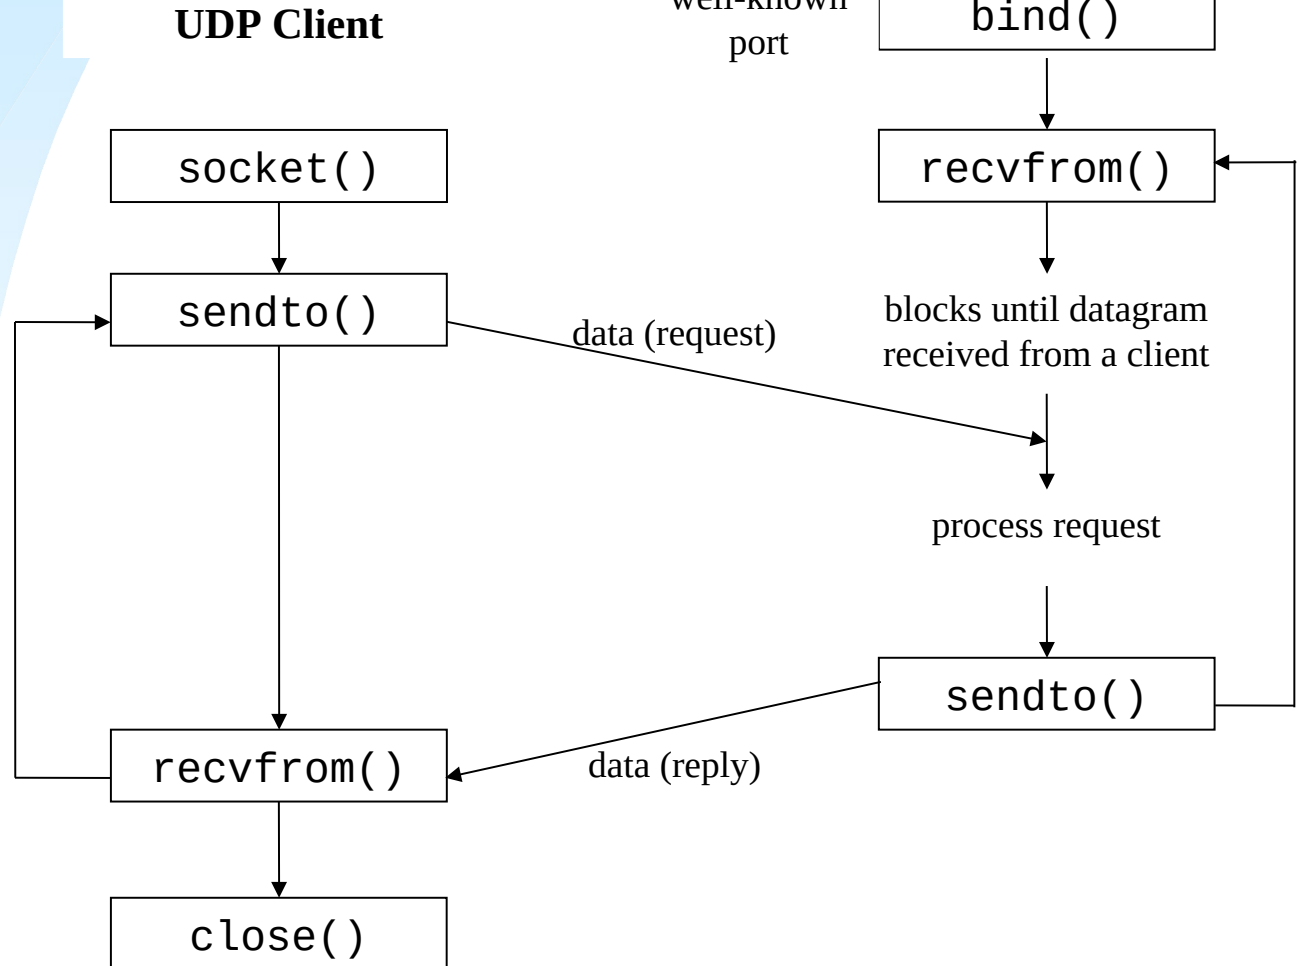
```
int recvfrom(int sockfd, char *buff, int nbytes, int flags, struct sockaddr *from, socklen_t *addrlen);
```

- Fills in the socket address structure pointed to by *from* with the protocol address of who sent the datagram
- If no incoming data is available at the socket, the recvfrom function blocks and waits for data to arrive according to the blocking rules

- *User datagram format :*



UDP Client-Server



Header Files

- `#include <sys/socket.h> // Core BSD socket functions and data structures.`
- `#include <netinet/in.h> // AF_INET and AF_INET6 address families and their corresponding protocol families PF_INET and PF_INET6.`
- `#include <arpa/inet.h> // Functions for manipulating numeric IP addresses.`
- `#include <netdb.h> // Name resolution`

Example (UDP_CLIENT)

- struct MSG
{
 int m_txt[MAX_TEXT];
 int m_seq;
 int total_no_of_packets;
};
- struct MSG msg;
- struct sockaddr_in sa;
- sock1 = socket(AF_INET,SOCK_DGRAM,0);
- bytes_sent = sendto(sock1, (void*)&msg, sizeof(struct MSG),0,(struct sockaddr*) &sa, sizeof(struct sockaddr_in));

Example (UDP_SERVER)

- struct MSG
{
 int m_txt[MAX_TEXT];
 int m_seq;
 int total_no_of_packets;
};
- struct MSG msg;
- int s_length;
- struct sockaddr_in sa,cli;
- sock1 = socket(AF_INET,SOCK_DGRAM,0);

Example UDP_SERVER

- `sa.sin_family = AF_INET;`
- `sa.sin_port = htons(PORT);`
- `sa.sin_addr.s_addr = htonl(INADDR_ANY);`
- `s_length=sizeof(sa);`
- `bind(sock1,(struct sockaddr *)&sa, s_length);`
- `rec=recvfrom(sock1,(void*)&msg,sizeof(struct MSG),0,(struct sockaddr *)&cli,sizeof(cli));`




TCP Protocol



Source port address 16 bits				Destination port address 16 bits				
Sequence number 32 bits								
Acknowledgment number 32 bits								
HLEN 4 bits	Reserved 6 bits	URG	ACK	PUSH	RESET	SYN	FIN	Window size 16 bits
Checksum 16 bits				Urgent pointer 16 bits				
Options and Padding								

THREE-WAY HANDSHAKING

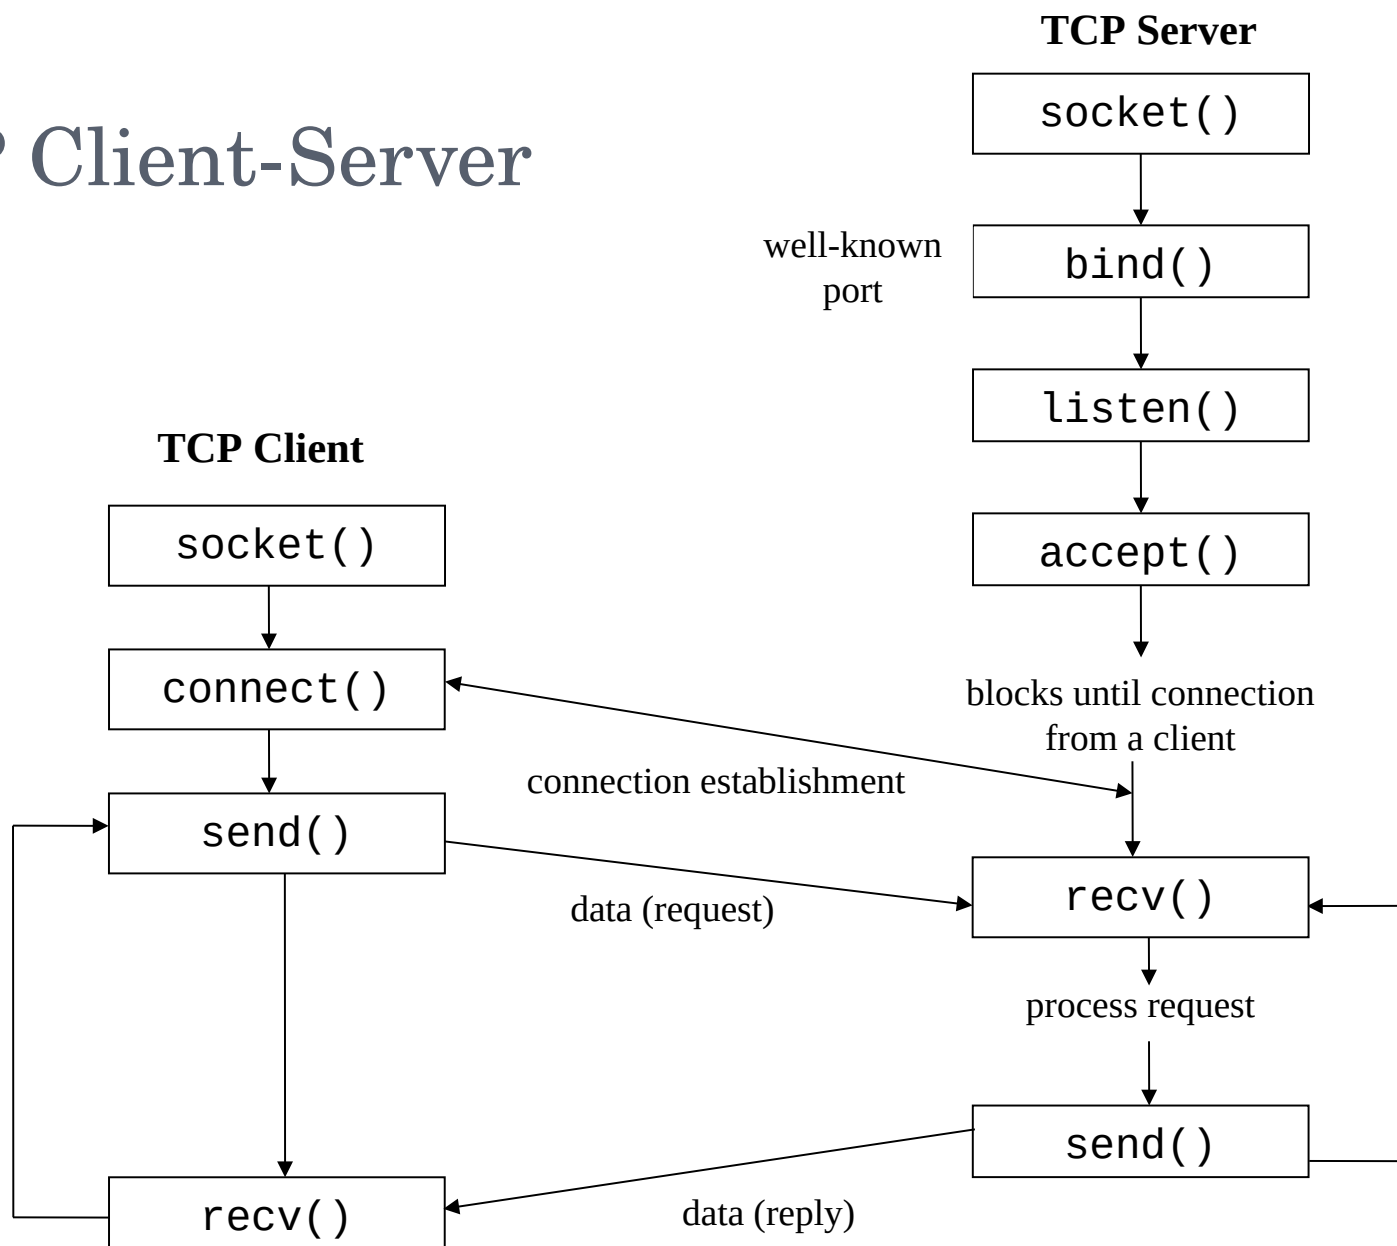


Passive Open: The application on the server is passive. The application is listening, awaiting a connection

The application on the client makes a connection request to the server where the application is passive open.

The application on the client is said to be "**active open**".

TCP Client-Server



connect System Call

```
int connect(int sockfd, struct sockaddr *servaddr,  
            int addrlen);
```

- For connection oriented protocols the connect system call results in the actual establishment of a connection between the local and foreign system
- *sockfd* is a socket descriptor that was returned by the socket system call
- *servaddr* is a pointer to an address structure containing the address of the server
- *addrlen* is the length of the address structure

connect()

```
#include <sys/socket.h>
int connect( int sockfd, const struct sockaddr * servaddr, socklen_t
    addrlen );
```

Returns: 0 if OK, -1 on error

- e.g. connect()
 - struct sockaddr_in servaddr;
 - ...
 - servaddr.sin_family = AF_INET;
 - servaddr.sin_addr.s_addr = htonl("220.95.133.100");
 - servaddr.sin_port = htons(7); // 7 for echo service

 - if(**connect**(*sockfd*, (struct sockaddr *)&servaddr,
sizeof(servaddr)) < 0)
 - { ... }
 - ...

listen System Call

```
int listen(int sockfd, int backlog);
```

- This system call is used by a connection-oriented server to indicate that it is willing to receive connection
- Called after both the `socket` and `bind` system calls, and immediately before the `accept` system call
- The *backlog* argument specifies **how many connection requests can be queued by the system** while it waits for the server to execute the `accept` system call
- That number is only the **size of the connection queue**, where new connections wait for somebody to accept them. As soon as your application calls `accept()`, a waiting connection is removed from that queue. So, you can definitely handle more than 128 simultaneous connections because they usually only spend a short time in the queue


Listen()




Queue

Accept()

As soon as your application calls `accept()`, a waiting connection is removed from that queue.



recv SYN send SYN+ACK



SYN Queue

recv ACK

Application



`accept()`



Accept Queue

Read/write
operation

Buffer

accept System Call

```
int accept(int sockfd, struct sockaddr *peer, int  
          *addr len);
```

- Takes first connection request on the queue and creates another socket with the same property as *sockfd*
- Blocks the caller if no connection request is pending
- *peer* argument return the address of the connected peer client
- *addr len* returns the length of the client address structure

listen() and accept()

```
#include <sys/socket.h>
int listen( int sockfd, int backlog );
```

Returns: 0 if OK, -1 on error

```
int accept( int sockfd, struct sockaddr * cliaddr, socklen_t *
addrlen )
```

Returns: non-negative descriptor if OK, -1 on error

- e.g. listen()
 - #define LISTENQ 5
 - ...
 - if(**listen**(*mysock*, **LISTENQ**) < 0) { ... }
 - while(1){
 - clilen = sizeof(*cliaddr*);
 - if((*clisockfd* = **accept**(*mysock*, (struct sockaddr *)&*cliaddr*,
&*clilen*)) < 0) { ... }
 - ...
 - }

recv() and send()

```
#include <sys/types.h>
#include <sys/socket.h>
int recv( int sockfd, void * msg, size_t len, int flags );
int send( int sockfd, const void * msg, size_t len, int flags );
```

Both return: number of bytes read or written if OK, -1 on error

msg: message you want to send

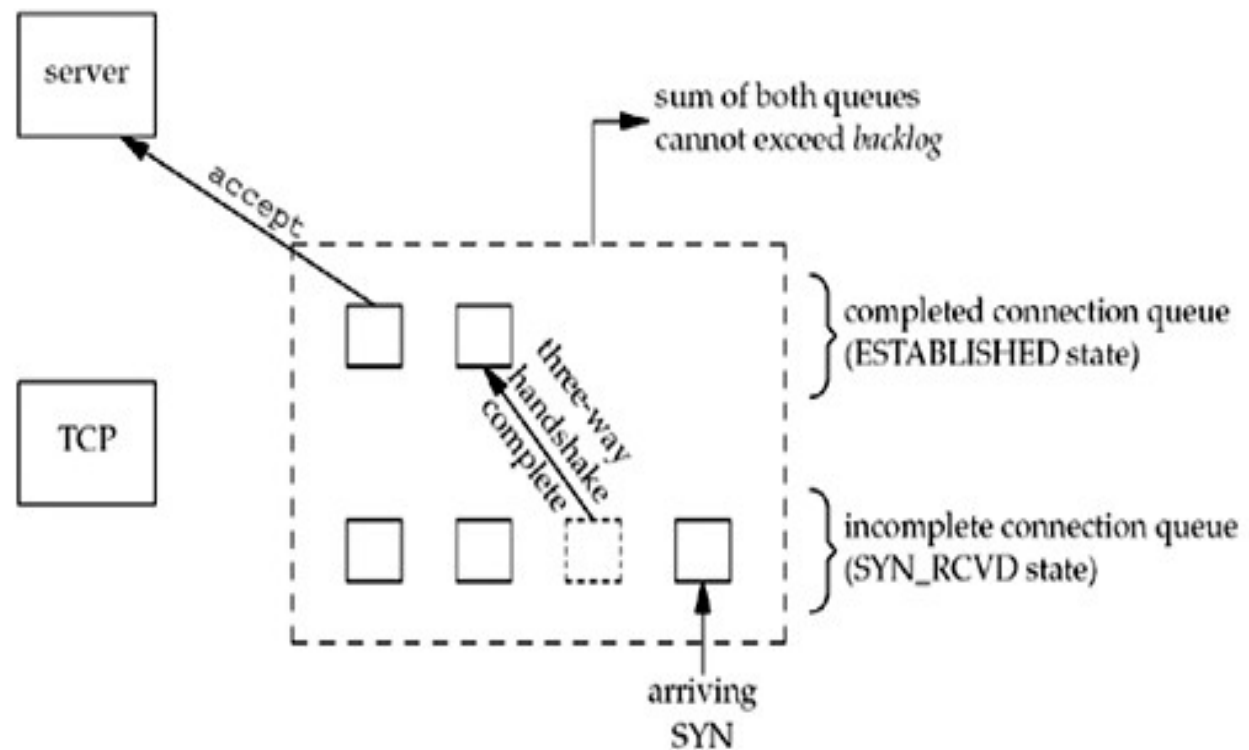
len: length of the message

flags := 0


returned: the number of bytes actually sent or received

- e.g. send() and recv()
 - if((n = send(**sockfd**, **buff**, **len**, 0)) != len){ ... }
 - ...
 - if((n = recv(**sockfd**, **buff**, **BUFMAX**, 0)) < 0) { ... }

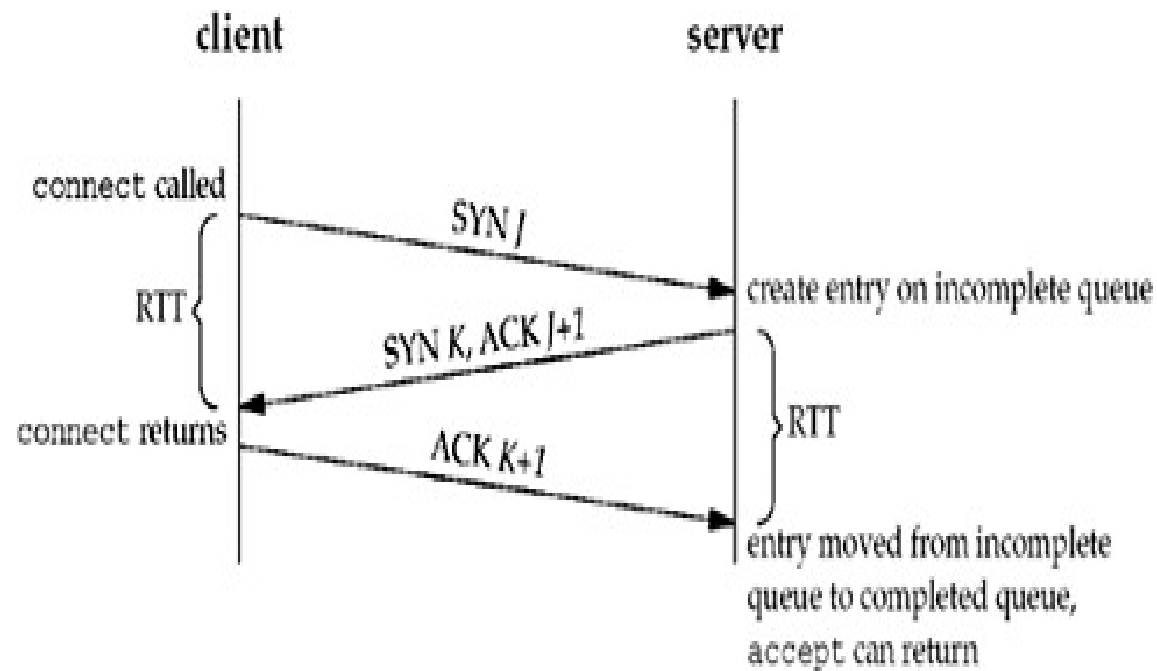
The two queues maintained by TCP for a listening socket.



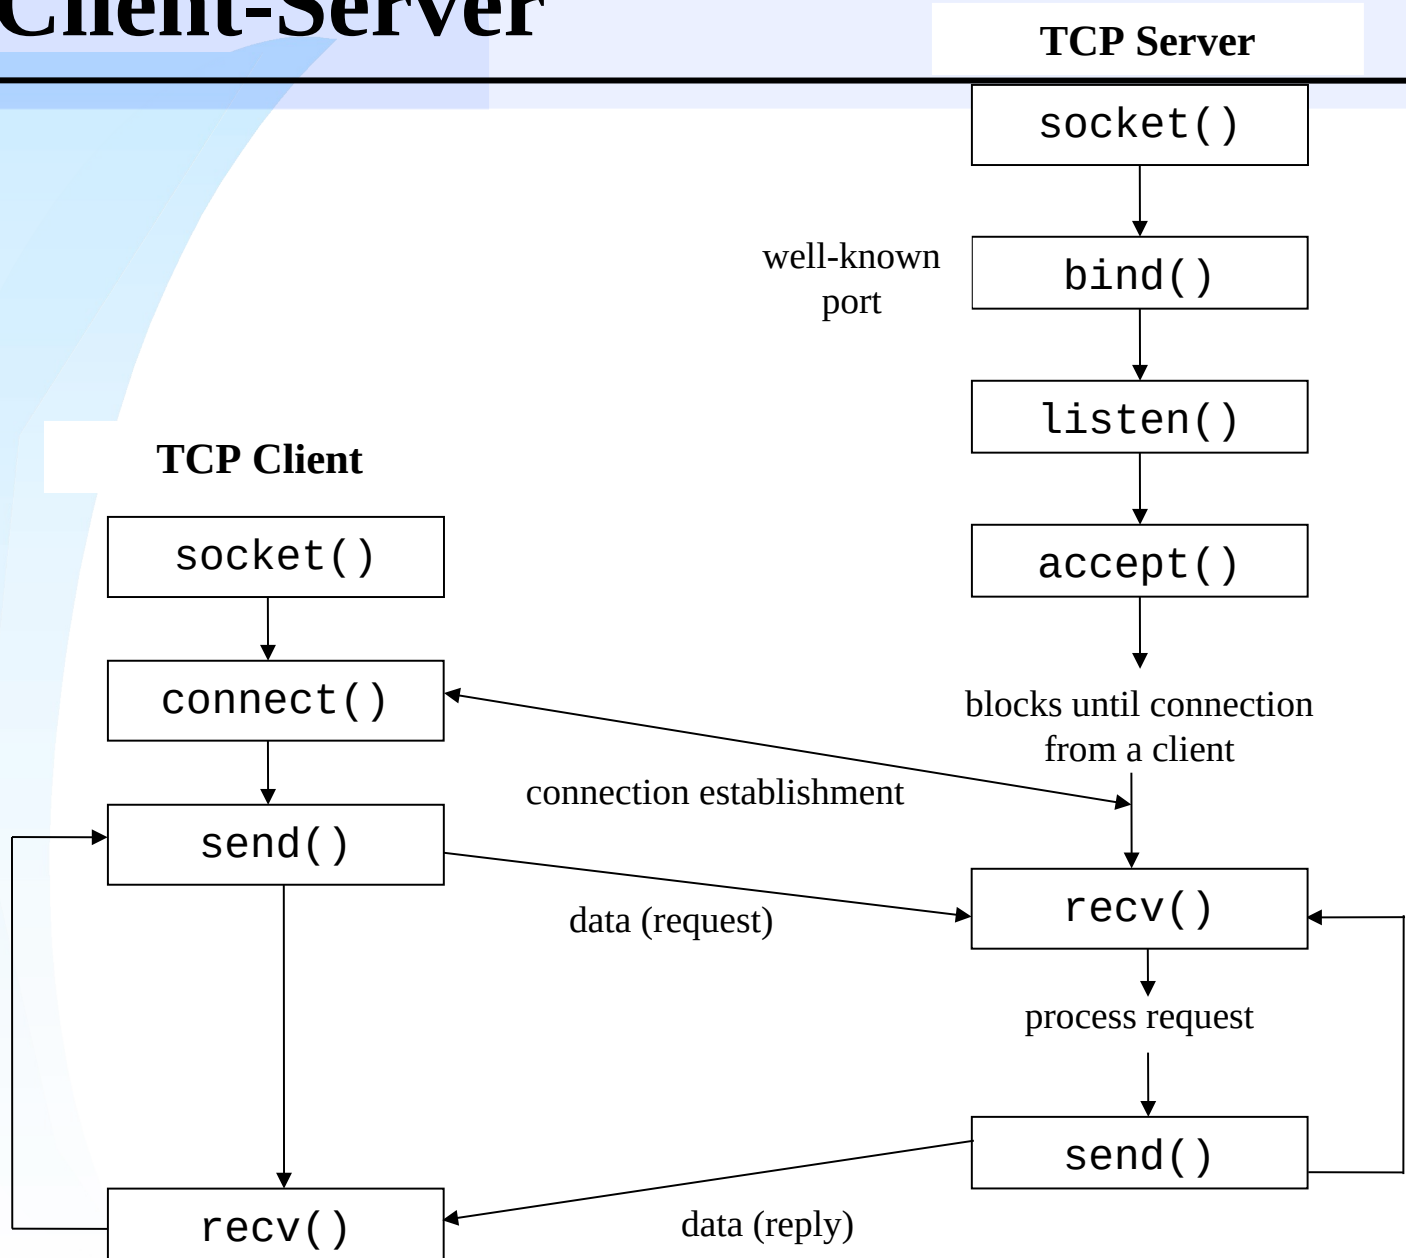
THREE-WAY HANDSHAKING



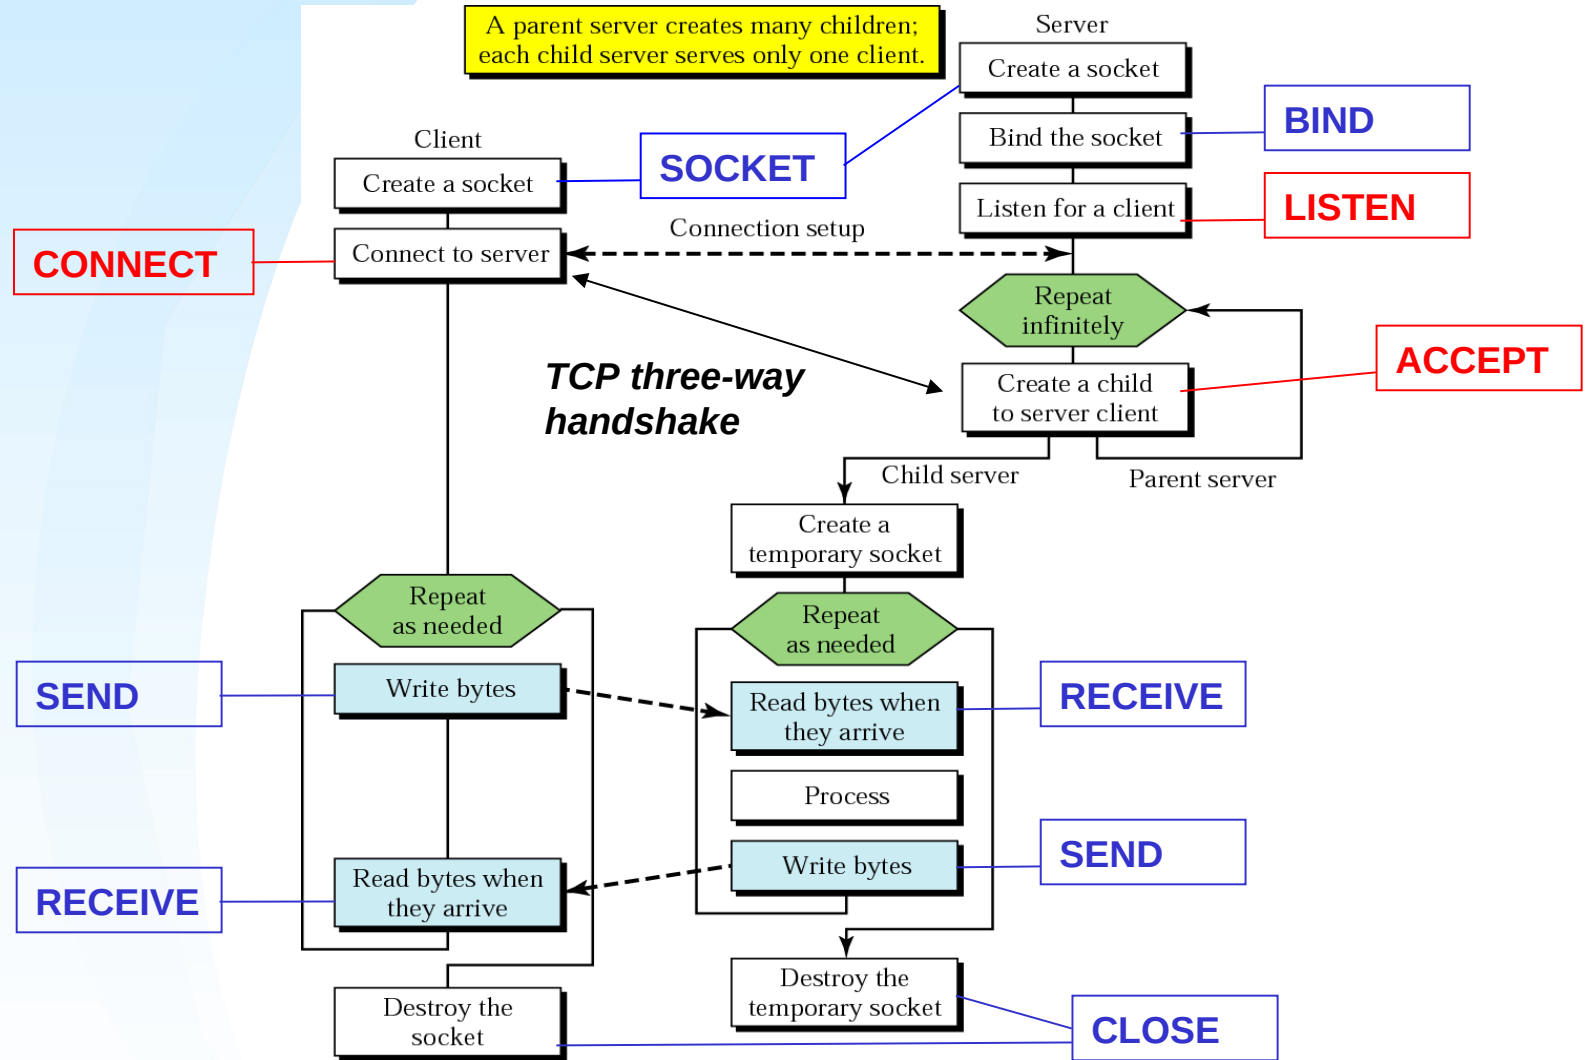
TCP three-way handshake and the two queues for a listening socket.



TCP Client-Server



Client+server: connection-oriented



Concurrent server

TCP SERVER

- struct sockaddr_in server_addr, client_addr;
- int sin_size;
- if ((sock = socket(AF_INET, SOCK_STREAM, 0)) == -1)
 perror("Socket");
 exit(1);
}
- server_addr.sin_family = AF_INET;
- server_addr.sin_port = htons(5000);
- server_addr.sin_addr.s_addr = INADDR_ANY;
- bzero(&(server_addr.sin_zero), 8);
- if (bind(sock, (struct sockaddr *)&server_addr, sizeof(struct sockaddr)) == -1)
{
 perror("Unable to bind");
 exit(1);
}

TCP SERVER

- ```
if (listen(sock, 5) == -1)
{
 perror("Listen");
 exit(1);
}
```
- ```
sin_size = sizeof(struct sockaddr_in);
```
- ```
connected = accept(sock, (struct sockaddr *)
&client_addr, &sin_size);
```
- ```
char send_data[1024],recv_data[1024];
```
- ```
gets(send_data);
```
- ```
send(connected, send_data,strlen(send_data), 0);
```

TCP CLIENT

```
■ struct sockaddr_in server_addr;

■ if ((sock = socket(AF_INET, SOCK_STREAM, 0)) == -1)
{
    perror("Socket");
    exit(1);
}

■ server_addr.sin_family = AF_INET;
■ server_addr.sin_port = htons(5000);
■ server_addr.sin_addr = inet_addr("server IP Address"); //inet_addr("192.168.148.12")
■ bzero(&(server_addr.sin_zero),8);

■ if (connect(sock, (struct sockaddr *)&server_addr, sizeof(struct sockaddr)) == -1)
{
    perror("Connect");
    exit(1);
}
```

TCP CLIENT

- `char recv_data[1024];`
- `bytes_recieved=recv(sock,recv_data,1024,0);`

References

- Unix Network Programming, Volume 1, Second Edition
 - W Richard Stevens
 - Pearson Education



